


How Sibr works

A complete breakdown: the big picture, one evaluation step by step, every parameter, audiences, security, logs and training. Updated 8 October 2026.

The name. Sibr is سبر (seen, baa, raa): to probe something and test its depth, as you would sound a well or a wound. Sibr probes a business idea to see how deep it really goes before you spend months building it. It is not صبر (sabr, patience).

In one paragraph. You describe a business idea. Sibr scores it on six parameters (pain, value, target audience, growth, Islamic alignment and competition), combines them into one Sibr score out of 100, and then works out which audiences would actually want it and pay for it. It is built for Muslim founders, so it checks whether the *business* is permissible. The customers can be anyone. There are two engines: **Quick** runs instantly in any browser with keyword rules, and **Laya** runs the real model on your PC. Everything is free and runs on your own machine.

1. The big picture




- **You and your peers** reach the PC only through Tailscale. Tailscale passes the Sibr website and API through on HTTPS port 443. The service itself listens only on 127.0.0.1, so nothing else on the internet can reach it.
- **Anyone** can open the public site and use the Quick engine, which runs entirely in their browser. Choosing "Laya on my PC" there sends the idea to your PC, and that only

works for devices on your Tailscale network.

- **On the PC**, one Python service does everything: it serves the website, checks the access key, runs the pipeline, talks to Ollama and Laya on the GPU, searches DuckDuckGo, and writes every request to the log.

2. One evaluation, step by step



Typical total: 30–75 s with the GPU free.
A game holding the GPU can make it time out.

Islamic alignment uses two checks: Laya's verdict and a keyword scan (riba, gambling, alcohol, speculation and similar). The stricter of the two is kept. Any audience cell whose Laya confidence is below 0.5 is marked "check" for a human look.

3. The six parameters

Parameter	What it asks	Levels	Weight
Pain	How painful is the problem for the customer?	0 no real pain · 1 nice to have · 2 painful, urgent	65% together (their average)
Value	How willing are customers to pay?	0 unlikely · 1 a little · 2 clearly willing	
Target audience	How clearly defined is the customer?	0 vague · 1 somewhat · 2 sharply defined	
Growth	How fast is the market growing?	0 shrinking/flat · 1 steady · 2 fast-growing	
Competition	How crowded is the space?	OPEN (1.0) · DIFF, room to stand out (0.65) · CROWDED (0.25)	20%
Islamic alignment	Can a Muslim founder run this business?	YES (1.0) · YES*, depends on content (0.7) · REVIEW (0.45) · FLAG (0)	15%

Laya gives each score as an expected level between 0 and 2, with probabilities for each level, which the report shows as bars. The level is divided by 2 to get 0–1. In the matrix, 0.67 or above shows as ✓, 0.33 or above as ~, and lower as ✕.

Worked example. Pain 1.8, value 1.6, target 1.2, growth 1.8 → average of (0.9, 0.8, 0.6, 0.9) = 0.80. Competition DIFF = 0.65. Islamic YES = 1.0.

Sibr = $100 \times (0.65 \times 0.80 + 0.20 \times 0.65 + 0.15 \times 1.0) = 100 \times (0.52 + 0.13 + 0.15) = \mathbf{80}$, so the verdict is "Strong" (75+). The other bands are Promising 60+, Risky 45+ and Weak below that. Any FLAG makes the verdict "Flagged", whatever the score.

The Islamic result is a flag for scholarly review, never a ruling. It is about the business model, not about who buys from it.

4. Audiences

Open (default, no filter)	The LLM suggests about 10 likely customer audiences for this exact idea, from any demographic, profession, business or interest. Laya scores each one. If the LLM is down, Sibr scores a broad mix of 25 segments instead. The Quick engine scans all 141 and shows the best 15.
Quick picks	Core 21 · Everyday consumers · Professionals · Businesses · Young people · Families · Muslim market · Everything (141).
Fine-tune	141 segments in 10 folded groups: geography (23), age (7), gender (2), income (4), life stage (12), profession (18), businesses (23), faith & culture (11), interest (33) and buyer type (8). There is a search box, and each group has All and None buttons.
Your own	Type up to 12 audiences, for example "dentists in Texas". Laya asks them the same question as catalogue segments: <i>"Would <audience> want this and be willing to pay for it?"</i>
Cost	Every audience adds one Laya question. 22 core questions plus up to 42 audiences fit in one pass; more audiences run in extra passes of 64. The audience box shows the count before you run.

5. Quick engine vs Laya

	Quick (any browser)	Laya on my PC
How	Keyword themes and fixed rules in engine.js	LLM idea card + web evidence + Laya typed verdicts + LLM write-up
Speed	Under a second	About 30–75 seconds
Needs	Nothing	Your PC on, Sibr running, Tailscale, the access key
Trust	Directional only	Calibrated probabilities and confidence, with evidence links

If Laya can't be reached, the site falls back to Quick and says why in the report.

6. Security: who can reach what

- **The service listens on 127.0.0.1 only.** Other devices reach it only through `tailscale serve` on port 443, inside your private Tailscale network. Funnel (the public internet) is never used.
- **Access key** (`SIBR_API_KEY`): required for every evaluation. Peers get this key.
- **Admin key** (`SIBR_ADMIN_KEY`): required for the logs. It is yours alone, and the access key cannot open the logs.

- **A Tailscale share exposes the whole PC by default.** Use the access-control rule in your guide so shared people can reach only port 443 on this PC.
- **Ollama listens on every network** (port 11434). Turn off "Expose Ollama to the network" in its settings. Sibr talks to it on the PC itself, so it keeps working.
- Both keys live in `server\.env`, which is never committed to GitHub. The site only allows its own files to be downloaded.


7. Live request logs

Open `/logs` on your PC or over Tailscale. Your guide has the full link with the admin key. The page refreshes every 2 seconds. Each row shows:

- the time
- where the request came from: this PC, or the Tailscale device and user
- the method and path
- the status (401 means a wrong key)
- how long it took
- for evaluations: the start of the idea, the audience mode and the Sibr score

The last 500 requests are kept in memory. The full history is in `server\logs\requests.jsonl`.

8. Training Laya on your own grades



The teacher only labels the audiences it was asked about. An audience it never saw is left out of training, not counted as "no". Older records cover the core 21.

9. What affects speed

- **The GPU being busy** matters most. A game can starve Laya and Ollama until requests time out.
- **Open mode** adds one LLM call to suggest audiences.
- **More audiences** mean more Laya questions and passes, and longer persona prompts.
- **Web evidence** depends on DuckDuckGo. Set `SIBR_EVIDENCE=off` to skip it.
- **Quick** is always instant.

10. Where things live

Website	<code>index.html</code> , <code>app.js</code> , <code>engine.js</code> , <code>segments.js</code> , <code>styles.css</code>
Service	<code>server/sibr_api/</code> : <code>app.py</code> (routes, keys, logs), <code>pipeline.py</code> (steps 2–7), <code>llm.py</code> , <code>laya_client.py</code> , <code>evidence.py</code> , <code>segments.py</code> + <code>segments.json</code> (audiences)
Training	<code>server/training/</code> : <code>make_dataset.py</code> , <code>finetune.py</code> , <code>data/</code> (<code>review.csv</code>)
Windows scripts	<code>scripts/</code> : <code>setup.ps1</code> , <code>start.ps1</code> , <code>tailscale.ps1</code> , <code>autostart.ps1</code>
Settings and keys	<code>server\.env</code> (private, never committed)

Sibr · سبر · github.com/Munazzar/sibr