


How Sibr works

Sibr (سبر, seen-baa-raa): to probe the depth of something. Not صبر (sabr, patience).
A complete breakdown of what runs where, what happens to an idea, how the score is made, and how requests reach your PC. Everything runs on your own computer; nothing is paid.

1. High-level snapshot




Your PC runs three things: **Ollama** (the free local LLM), the **Sibr service** (which also holds the **Laya** model in memory), and **Tailscale**. The Sibr service listens only on 127.0.0.1, so nothing on your home network or the internet can reach it directly. Your devices reach it only through Tailscale, which encrypts traffic and only admits devices signed in to your tailnet or ones you've shared with. The public site works without any of this using the simpler Quick engine.

Component	What it does
Website (index.html, app.js, engine.js)	The page you use. Served by your PC at the ts.net address, or by GitHub Pages publicly. The Quick engine is a keyword-based estimate that runs in the browser; "Laya on my PC" sends the idea to your PC.
Sibr service	Python FastAPI app (server\sibr_api). Checks keys, logs requests, runs the pipeline, serves the site.
Ollama + llama3.2:3b	A 3-billion-parameter open model running on your NVIDIA GPU. Writes the idea card and the plain-English explanations.
Laya	A model built for answering typed questions (scores, choices, yes/no) with calibrated confidence. Gives the actual verdicts. Can be fine-tuned on your own grades.
DuckDuckGo search	Free web search for competitors and discussions; no key or account.

Tailscale	Private encrypted network between your devices. <code>tailscale serve</code> publishes Sibr on HTTPS inside the tailnet only (never funnel).
-----------	--

2. What happens when you press Run (detailed)



3. The parameters and how they're scored

Parameter	Question Laya answers · levels
Pain	How painful is the problem this idea solves for its customer? · no real pain / nice to have / painful, urgent problem
Value	How willing are customers to pay for this? · unlikely to pay / might pay a little / clearly willing to pay
Target	How clearly defined is the target customer? · vague or everyone / somewhat defined / sharply defined customer
Growth	How fast is the market for this growing? · shrinking or flat / steady / fast-growing
Competition	OPEN (few or no direct competitors, 1.0) · DIFF (competitors exist but clear differentiator, 0.65) · CROWDED (many similar products, 0.25)

Islamic alignment	YES (permissible, 1.0) · YES* (depends on content or advertising, 0.7) · REVIEW (may involve interest or speculation, needs scholarly review, 0.45) · FLAG (gambling, alcohol or clearly prohibited, 0). A keyword check of your text can make this stricter but never looser.
Audience fit	141 audience segments in 10 groups: Geography (23), Age group (7), Gender (2), Income (4), Life stage (12), Profession (18), Businesses (23), Faith & culture (11), Interest (33), Buyer type (8). Customers do not need to be Muslim; the Islamic check is about the product itself, not who buys it. Each scored audience gets a 0–100% chance that it wants the product and pays for it.
Audience modes	Pick segments or presets (only those are scored, so fewer is faster). Open (no filter): the LLM proposes the 10 most likely customer audiences for your idea and Laya scores them; without a working LLM a broad mix of 25 is used. Your own : type up to 12 audiences such as "dentists in Texas". With nothing chosen, the core 21 are scored.
Confidence	Laya reports how sure it is. Cells below 50% are marked for a human check rather than trusted.

Worked example (halal investing app for students): pain 0.49, value 0.81, target 0.79, growth 0.77 → average 0.72; competition OPEN = 1.0; Islamic REVIEW = 0.45. Sibr = $100 \times (0.65 \times 0.72 + 0.2 \times 1.0 + 0.15 \times 0.45) \approx 73$.


4. Security: who can reach what

Layer	Protection
Network	Sibr listens on 127.0.0.1 only. No router or Windows Firewall ports were opened. Remote access exists only through Tailscale serve, inside your tailnet. Funnel (public internet) is never used.
Access key	/evaluate needs the 24-character access key (X-Sibr-Key). Wrong key → 401, logged.
Admin key	/logs needs a separate admin key that peers don't get.
Files	The service serves only the four site files; server\.env and everything else in the folder are unreachable.
Shared peers	A Tailscale share reaches the whole PC by default (for example Ollama on port 11434 and Windows file sharing, which still needs your Windows password). Restrict shared users to port 443 in the Tailscale access controls and turn off "Expose Ollama to the network" before sharing widely.

5. Live request logs

Open `http://localhost:8787/logs` on the PC or `https://your-pc.your-tailnet.ts.net/logs` from your devices, and enter the admin key once (it's in your Sibr guide and in `server\.env` as `SIBR_ADMIN_KEY`). The page refreshes every 2 seconds and shows each request: time, where it came from (this PC, or the Tailscale device or user), what was asked, the result code, how long it took, and for evaluations the start of the idea and its score. Filters show evaluations only, Tailscale only, or errors and denied requests. Everything is also saved to `server\logs\requests.jsonl`, so history survives restarts.

6. Training Laya on your own grades



Stock Laya answers "zero-shot" (it has never seen Sibr's questions). Its authors report about 0.36 accuracy before fine-tuning and 0.77 after on their benchmark. The local LLM drafts a grade for each of 1,000 sample ideas (about 26 seconds each), you correct the grades you disagree with, and the GPU fine-tune teaches Laya your judgement. Then Sibr loads the new checkpoint instead of the stock one.

7. Speed

A typical evaluation takes 30–75 seconds, spread over the two LLM calls, the web searches and the Laya pass; each result shows exactly how long every step took. Scoring fewer audiences shortens it; Open mode adds one extra LLM call to propose audiences. Games or other heavy GPU work on the PC slow everything down sharply, because Ollama, Laya and the game share the same graphics card.